

Towards a Converged Relational-Graph Optimization Framework

Anonymous Authors

ABSTRACT

The recent ISO SQL:2023 standard adopts SQL/PGQ (Property Graph Queries), facilitating graph-like querying within relational databases. This advancement, however, underscores a significant gap in how to effectively optimize SQL/PGQ queries within relational database systems. To address this gap, we extend the foundational SPJ (Select-Project-Join) queries to SPJM queries, which include an additional matching operator for representing graph pattern matching in SQL/PGQ. Although SPJM queries can be converted to SPJ queries and optimized using existing relational query optimizers, our analysis shows that such a graph-agnostic method fails to benefit from graph-specific optimization techniques found in the literature. To address this issue, we develop a converged relational-graph optimization framework called RelGo for optimizing SPJM queries, leveraging joint efforts from both relational and graph query optimizations. Using DuckDB as the underlying relational execution engine, our experiments demonstrate that RelGo can generate efficient execution plans for SPJM queries. On well-established benchmarks, these plans exhibit an average speedup of 14.72× compared to those produced by the graph-agnostic optimizer.

1 INTRODUCTION

In the realms of data management and analytics, relational databases have long been the bedrock of structured data storage and retrieval, empowering a plethora of applications. The ubiquity of these databases has been supported by the advent of Structured Query Language (SQL) [8], a standardized language that has been adopted widely by various relational database management systems for managing data through schema-based operations.

Despite its considerable success and broad adoption, SQL has its limitations, particularly when it comes to representing and querying intricately linked data. Consider, for instance, the relational tables of Person and Knows, the latter symbolizing a many-to-many relationship between instances of the former. Constructing a SQL query to retrieve a group of four persons who are all mutually acquainted is not a straightforward endeavor, potentially leading to a cumbersome and complex SQL expression.

In comparison, such a scenario could be succinctly addressed using graph query languages, exemplified by Cypher [3], a language where queries are expressed as intuitive and concise graph pattern matching. This discrepancy between the relational and graph querying paradigms has given rise to the innovative SQL/Property Graph Queries (SQL/PGQ), an extension formally adopted in the ISO/IEC SQL:2023 standard [34]. SQL/PGQ is designed to amalgamate the extensive capabilities of SQL with the inherent benefits of graph pattern matching. With SQL/PGQ, it is now possible to define and query graphs within SQL expressions, transforming otherwise complex relational queries – characterized by multiple joins – into simpler and more intuitive graph queries.

EXAMPLE 1. Consider the four relational tables in the database: Person(id, name, place_id), Message(id, content, date),

Like(p_id, m_id, date), and Place(id, name). Using SQL/PGQ, a property graph G is articulated as a **GRAPH_TABLE**, established on the basis of the first three tables. In this mapping, rows from Person and Message are interpreted as vertices with labels “Person” and “Message” respectively, while rows from Like represent edges with the label “Likes”. This mapping process will be formally elaborated in Def. 1. An SQL/PGQ query to discover the friends of a person named “Tom” and the place they live in, where “Tom” and friends share an affinity for the same message, can be formulated as:

```
SELECT p2_name, place.name
FROM GRAPH_TABLE (G
  MATCH (p1:Person)-[:Likes]->(m:Message),
        (p2:Person)-[:Likes]->(m:Message),
        (p1)-[:Knows]->(p2)
  COLUMNS (
    p1.name AS p1_name,
    p1.place_id AS p1_place_id,
    p2.name AS p2_name)
) g
JOIN Place p ON g.p1_place_id = p.id
WHERE g.p1_name = 'Tom';
```

In graph G , a **GRAPH PATTERN MATCHING** is employed to decode the intricate relationships between persons and messages. Upon executing the pattern matching, a **COLUMNS** clause projects the results into a tabular format, enumerating essential attributes. Then the **RELATIONAL JOIN** is performed on resultant table g and Place table to obtain the place’s name.

The SQL/PGQ standard, while a significant leap forward in the realm of relational databases, primarily addresses language constructs. A discernible gap exists in the theoretical landscape, particularly in analyzing, transforming, and optimizing SQL/PGQ queries with hybrid relational and graph semantics.

Relational query optimization has historically leaned on the SPJ (selection-projection-join) skeleton [10, 37], which provides a systematic approach for analyzing query complexity [9, 19], devising heuristic optimization rules [11, 13], and computing optimal join order [12, 15]. Recently, graph techniques have been introduced to optimize relational queries [15, 21, 29, 30]. In particular, GrainDB [21] introduced a predefined join operator that materializes the adjacency list (rows) of vertices, enabling more efficient join execution. While these techniques can be empowered by graph techniques, they target purely relational query rather than the relational-graph hybrid query of SQL/PGQ.

In parallel to relational query optimization, significant strides have been made in optimizing graph pattern matching. A common practice is to leverage join-based techniques to optimize the query [4, 24, 25, 43]. Scalable join algorithms, such as binary-join [24], worst-case optimal join [4], and their hybrid variants [26, 31, 43], have been proposed for solving the problem over large-scale graphs. However, despite the effectiveness of these techniques for pattern matching on graphs, they cannot be directly applied to relational databases due to the inherent differences in data models.

In this paper, we propose the first converged optimization framework, RelGo, that optimizes relational-graph hybrid queries in a relational database, in response to the advent of SQL/PGQ. A straightforward implementation [1, 40, 41] can involve directly transforming the graph component in SQL/PGQ queries into relational operations, allowing the entire query to be optimized and executed in any existing relational engine. While we contribute to building the theory to make such a transformation workable, this *graph-agnostic* optimization approach suffers from several issues, including graph-unaware join orders, suboptimal join plans, and increased search space, as will be discussed in Sec. 3.1.2.

To address these challenges, RelGo is proposed to leverage the strengths of both relational and graph query optimization techniques. Building upon the foundation of SPJ queries, we introduce the SPJM query skeleton, which extends SPJ with a matching operator to represent graph queries. We adapt state-of-the-art graph optimization techniques, such as the decomposition method [43] and the cost-based optimizer [26], to the relational context, effectively producing worst-case optimal graph subplans for the matching operator. To facilitate efficient execution of the matching operator, we introduce graph index inspired by GrainDB's predefined join [21], based on which graph-based physical operations are implemented. The relational part of the query, together with the optimized graph subplans encapsulated within a special operator called SCAN_GRAPH_TABLE, is then optimized using standard relational optimizers. Finally, we incorporate heuristic rules, such as FilterIntoMatchRule, to handle cases unique to SPJM queries that involve the interplay between relational and graph components.

We have made the following contributions in this paper:

- (1) We introduce the concept of RGMMapping (*Relational-to-Graph Mapping*), which serves as the foundation for mapping relational data models to property graph models as specified by SQL/PGQ. Based on this concept, we define a new query skeleton called SPJM to better analyze the relational-graph hybrid queries. (Sec. 2)
- (2) We construct the theory for transforming any SPJM query into an SPJ query. Such a graph-agnostic approach enables existing relational databases to handle SPJM queries without low-level modifications. We also formally prove that the search space of the graph-agnostic approach is exponentially larger than our solution. (Sec. 3)
- (3) We introduce RelGo, a converged optimization framework that leverages the strengths of both relational and graph query optimization techniques to optimize SPJM queries. RelGo adapts state-of-the-art graph optimization techniques to the relational context, and implements graph-based physical operations based on graph index for efficient query execution. (Sec. 4)
- (4) We develop RelGo by integrating it with the industrial relational optimization framework, Calcite [14], and employing DuckDB [2] for execution runtime. We conducted extensive experiments to evaluate its performance. The results on the LDBC Social Network Benchmark [27] indicate that RelGo significantly surpasses the performance of the graph-agnostic baseline, with an average speedup of 14.72 \times , and 10 \times even after graph index is enabled for the baseline. (Sec. 5)

This paper is organized in the order of the contributions. We survey related work in Sec. 6 and conclude the paper in Sec. 7.

2 PRELIMINARIES

2.1 Data Model

A schema, denoted as $S = (a_1, a_2, \dots, a_n)$, is a collection of attributes. Each attribute a_i in this schema is associated with a specific data domain D_i , which defines the set of permissible values that a_i can take. A relation R is defined as a set of tuples. We consider R to be a relation over schema S , if and only if, every tuple $\tau = (d_1, d_2, \dots, d_n)$ in R adheres to the schema's constraints, such that the value d_i for each position in the tuple corresponds to the data domain D_i of the attribute a_i in S . In other words, each value d_i in a tuple τ is drawn from the appropriate data domain D_i for its corresponding attribute a_i . Moreover, for any tuple τ in the relation R , the notation $\tau.a_i = d_i$ signifies that the attribute a_i in tuple τ has the value d_i . A table is a representation of a relation, where each row corresponds to a tuple in the relation, and each column represents an attribute in the schema. In this paper, we use the terms of relation and table interchangeably.

We define a *Property Graph* as $G = (V_G, E_G)$, where

- V stands for the set of vertices.
- Let $E \subseteq V \times V$ denote the set of edges in the graph. An edge $e \in E$ is represented as an ordered pair $e = (v_s, v_t)$, where $v_s \in V$ is the source vertex and $v_t \in V$ is the target vertex, indicating that the edge e connects from v_s to v_t .
- For any graph element ϵ that is either a vertex or an edge, we denote $\text{id}(\epsilon)$ and $\ell(\epsilon)$ as the globally unique ID and the label of ϵ , respectively. Given an attribute a_i , $\epsilon.a_i$ denotes the value of the attribute a_i of ϵ .

Given a vertex v , we denote its adjacent edges as $\mathcal{N}_G^E(v) = \{e = (v, v_t) | e \in E\}$ and its adjacent vertices (i.e., neighbors) as $\mathcal{N}_G(v) = \{v_t | (v, v_t) \in E\}$. It is important to note that the adjacent edges and vertices can be defined for both directions of an edge $e = (v_s, v_t)$, i.e., when $v = v_s$ or $v = v_t$. However, for simplicity, we only define one direction in this notation. In the actual semantics of the paper, both directions may be considered. The degree of v is defined as $d_G(v) = |\mathcal{N}_G(v)|$, and the average degree of all vertices in the graph is $\bar{d}_G = \frac{1}{|V_G|} \sum_{v \in V_G} d_G(v)$. In the rest of the paper, when the context is clear, we may remove G from the subscript for simplicity, for example $G = (V, E)$.

Considering two graphs G_1 and G_2 , we assert that G_2 is a subgraph of G_1 , symbolized as $G_2 \subseteq G_1$, if and only if $V_{G_2} \subseteq V_{G_1}$, and $E_{G_2} \subseteq E_{G_1}$. Furthermore, G_2 qualifies as an induced subgraph of G_1 under the condition that G_2 is already a subgraph of G_1 , and for every pair of vertices in G_2 , any edge e that exists between them in G_1 must also present in G_2 .

To formalize the integration of graph syntax within the realm of relational data, we introduce the concept of a *Relations-to-Graph Mapping* (i.e. RGMMapping), to facilitate the transformation of relational data structures into a property graph.

DEFINITION 1 (RGMMapping). *Given two sets of relations $\{R_{p_1}, \dots, R_{p_n}\}$ and $\{R_{q_1}, \dots, R_{q_m}\}$, we define an RGMMapping to map the relations to a property graph $G = (V, E)$, which are elaborated as follows:*

- **Vertex Mappings:** $\zeta_v : \bigcup_{i=1}^n R_{p_i} \mapsto V$ is a bijective function that maps every tuple $\tau \in R_{p_i}$ for all $1 \leq i \leq n$ to a unique vertex $v \in V$. This vertex v is assigned: an ID $\text{id}(v)$, a label $\ell(v)$ that matches the name of R_{p_i} , and attributes $v.\text{attr}^*$ that mirror the

attributes attr^* of τ . The relations that are mapped to vertices in the graph are referred to as vertex relations.

- **Edge Mappings:** $\zeta_e : \bigcup_{i=1}^m R_{q_i} \mapsto E$ is a bijective function that maps every tuple $\tau \in R_{q_i}$ for all $1 \leq i \leq m$ to a unique edge $e = (v_s, v_t) \in E$. Similar to vertices, each edge e is assigned: an ID $\text{id}(e)$, a label $\ell(e)$ corresponding to the name of R_{q_i} , and attributes $e.\text{attr}^*$ that reflect the attributes $\tau.\text{attr}^*$ of τ . For each edge relation R_{q_i} , there are two corresponding vertex relations R_{p_s} and R_{p_t} . We further define two total functions: $\lambda_{q_i}^s : R_{q_i} \rightarrow R_{p_s}$ and $\lambda_{q_i}^t : R_{q_i} \rightarrow R_{p_t}$, which maps a tuple $\tau \in R_{q_i}$ to a tuple $\tau_s \in R_{p_s}$ and a tuple $\tau_t \in R_{p_t}$, respectively. The vertices v_s and v_t corresponding to τ_s and τ_t , respectively, are obtained by applying the vertex mapping function ζ_v , i.e., $v_s = \zeta_v(\tau_s)$ and $v_t = \zeta_v(\tau_t)$. The relations R_{q_i} that are mapped to edges in the graph are referred to as edge relations.

EXAMPLE 2. An RGMMapping can be defined following the grammar of SQL/PGQ with **CREATE PROPERTY GRAPH** statements. Given the relations, an RGMMapping is defined with the statement shown in Fig. 1(a). The described RGMMapping involves assigning tuples from vertex relations, such as R_{Person} and R_{Message} , to graph vertices. For instance, the vertex v_{p_1} is associated with the tuple τ_{p_1} in R_{Person} , and thus assigned the label “Person” and the name attribute “Tom”. Similarly, edge relations R_{Likes} and R_{Knows} correspond to graph edges. To further elaborate on the mapping process, especially for the source and target vertex mappings of an edge, let’s consider the edge e_{l_1} . This edge originates from the tuple τ_{l_1} in the R_{Likes} relation. The source vertex v_{p_1} of e_{l_1} is linked to the tuple τ_{p_1} in R_{Person} via the function λ_{Likes}^s , following the condition “ $\tau_{l_1}.\text{pid} = \tau_{p_1}.\text{person_id}$ ”. Similarly, the target vertex v_{m_1} is associated with the tuple τ_{m_1} in R_{Message} via the function λ_{Likes}^t , following the condition “ $\tau_{l_1}.\text{mid} = \tau_{m_1}.\text{message_id}$ ”. As a result of this mapping, the edge e_{l_1} is assigned the label “Likes” and the attribute “date” with the value “2024-03-31”.

2.2 Matching Operator

Consider a property graph $G(V_G, E_G)$, alongside a *connected* pattern graph, represented as $\mathcal{P}(V_{\mathcal{P}}, E_{\mathcal{P}})$. Here, \mathcal{P} is a property graph that does not possess attributes, and we denote n and m as the number of vertices and edges in the \mathcal{P} , respectively. Graph pattern matching seeks to determine all subgraphs in G that are *homomorphic* to \mathcal{P} . Formally, given a subgraph $g \subseteq G$, a homomorphism from \mathcal{P} to g is a *surjective*, total mapping $f : V_{\mathcal{P}} \cup E_{\mathcal{P}} \rightarrow V_g \cup E_g$ that satisfies the following conditions: (1) For every vertex $u \in V_{\mathcal{P}}$, there is a corresponding vertex $v = f(u) \in V_g$ with $\ell(v) = \ell(u)$; (2) For each edge $e = (u_s, u_t) \in E_{\mathcal{P}}$, there is a corresponding edge $f(e) = (v_s, v_t) \in E_g$, ensuring that the mapping preserves the edge’s label, as well as its source and target vertices, that is $\ell(e) = \ell(f(e))$, and $f(u_s) = v_s$, $f(u_t) = v_t$. It’s important to highlight the homomorphism semantics, as one of the widely used semantics for graph pattern matching [5], do not require each pattern vertex and edge being uniquely mapped to distinct vertices and edges in the data graph. This facilitates a seamless integration between graph pattern matching and relational operations, but alternative semantics for graph pattern matching such as isomorphism can also be adopted, as will be further discussed in Sec. 3.1.

The outcomes of graph pattern matching can be succinctly modeled as a relation $GR_{\mathcal{P}}^G$, or more compactly $GR^{\mathcal{P}}$ in clear contexts,

defined over the schema $S = V_{\mathcal{P}} \cup E_{\mathcal{P}}$. Here, the sets V_G and E_G serve as the respective domains for the vertices and edges identified through the matching process. Within this framework, we refer to such a relation as a *Graph Relation*, a construct where all attributes are derived from the domain of a property graph. It is essential to recognize that any property graph G can be conceptualized as a graph relation GR^G , represented by a singular tuple that collectively encompasses all of its vertices and edges. Throughout this paper, we treat the notions of a property graph and a tuple of graph relation as essentially interchangeable terms. In alignment with this perspective, we elaborate on the *Matching* operator as follows.

DEFINITION 2 (MATCHING OPERATOR, \mathcal{M}). The *Matching Operator*, denoted as \mathcal{M} , is designed to perform graph pattern matching on a given graph relation GR against a specified pattern graph \mathcal{P} . For each graph instance g within GR , \mathcal{M} identifies all subgraphs of g that are homomorphic to \mathcal{P} , and subsequently, aggregates these mappings to construct a comprehensive graph relation. The operation of the matching Operator can be formally articulated as $\mathcal{M}(GR, \mathcal{P}) = \bigcup_{g \in GR} GR_g^{\mathcal{P}}$.

EXAMPLE 3. Let G denote the property graph derived from the relations via RGMMapping in Example 2. Given a pattern graph \mathcal{P} shown in Fig. 1(b), the results of graph pattern matching are subgraphs of G that are homomorphic to \mathcal{P} , which are represented as a graph relation $GR^{\mathcal{P}} = \mathcal{M}(GR^G, \mathcal{P})$, wherein each tuple corresponds to one of the matched subgraph.

This definition ensures that the matching operator is inherently closed regarding graph relations, which adheres to the language opportunities of “nested matching” (specified as PGQ-079) in SQL/PGQ [34]. In this paper, we only handle cases where G represents the entire property graph, and thereafter simplify the matching operator notation to $\mathcal{M}(\mathcal{P})$ when the context is clear.

2.3 Problem Definition

To study relational query optimization, it is common to focus on a category of queries known as SPJ queries, which encapsulate the three most frequently employed operations in database management: select, project, and (natural) join. These operations form the backbone of many relational queries. Given a set of relations R_1, R_2, \dots, R_m , an SPJ query is formally represented as:

$$Q = \pi_A(\sigma_{\Psi}(R_1 \bowtie \dots \bowtie R_m)).$$

Inspired from the SPJ paradigm, we introduce a novel category of queries, termed SPJM queries, to logically formulate SQL/PGQ [34] queries that blend relational and graph operations. Our focus is not on an exhaustive exploration of SQL/PGQ, which encompasses a broad spectrum of operations such as shortest-path, but rather on understanding and enhancing the core functional elements of such queries. The SPJM framework augments SPJ queries by incorporating a matching operator, thereby enriching the query’s expressive power to seamlessly navigate both relational and graph data domains. Given the set of relations and a property graph G constructed from these relations via an RGMMapping (Def. 1), an SPJM query is articulated as:

$$Q = \pi_A(\sigma_{\Psi}(R_1 \bowtie \dots \bowtie R_m \bowtie (\hat{\pi}_{A^*} \mathcal{M}_G(\mathcal{P})))) \quad (1)$$

In this formulation, $\hat{\pi}_{A^*} \mathcal{M}_G(\mathcal{P})$ is the *graph component* of the query, while the remaining part of the query is an SPJ expression

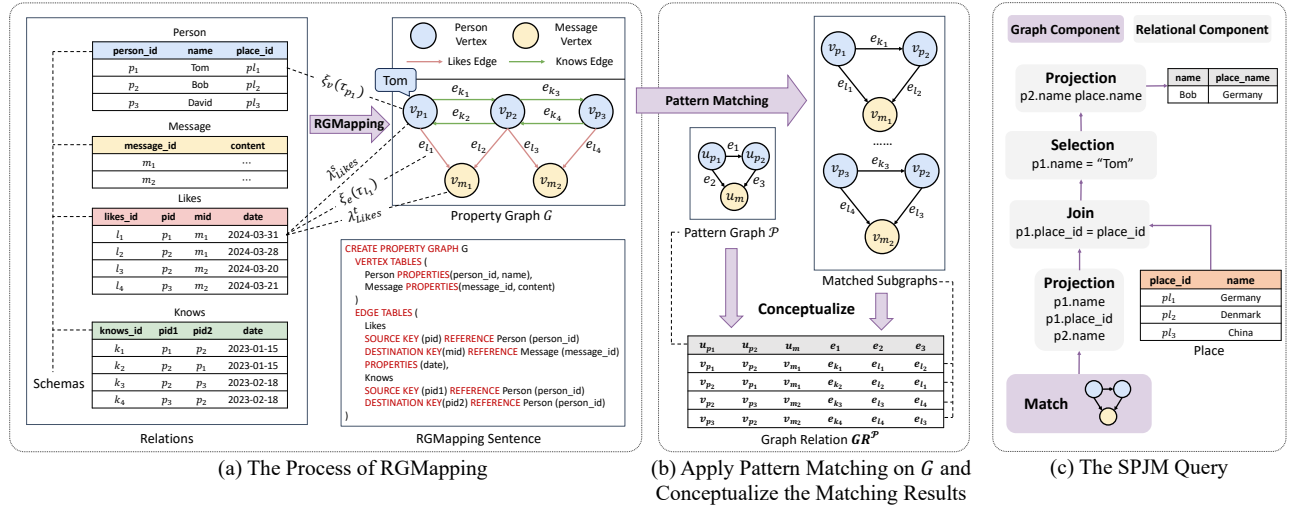


Figure 1: An example of RGMapping.

referred to as the *relational component*. Here, $M_G(\mathcal{P})$ represents the process of matching the pattern \mathcal{P} on the graph G and returns a graph relation as defined in Def. 2. The operator $\hat{\pi}_{A*}$ is a graph-calibrated projection operator that extracts the ID, label, and other attributes from the vertices and edges in the matched results. This process helps “flatten” graph elements into relational tuples. For example, given a graph relation GR that contains a vertex of $\{ID:0, \text{label:Person, name:“Tom”}\}$, the following projection

$$\hat{\pi}_{id(v) \rightarrow v_id, \ell(v) \rightarrow v_label, v.name \rightarrow v_name}(GR)$$

turns the vertex into a relational tuple of $(0, \text{Person, “Tom”})$. The projection operation is designed to reflect the **COLUMNS** clause in SQL/PGQ to retrieve specific attributes from vertices and edges as required. For simplicity, we assume that all attributes are extracted unless otherwise specified.

In this paper, we study the problem of optimizing SPJM queries in Eq. 1. Fig. 1(c) illustrates the SPJM query skeleton corresponding to the SQL/PGQ query in Example 1.

Table 1 summarizes frequently used notations in this paper.

3 OPTIMIZING MATCHING OPERATOR

In this section, we focus on handling the matching operator (more precisely, the graph component), which plays a distinct role within the SPJM queries compared to the SPJ queries. We discuss two main perspectives of optimizing the matching operator: logical transformation and physical implementation. Logical transformation is responsible for transforming a matching operator into a logically equivalent representation, while physical implementation focuses on how the matching operator can be efficiently executed.

3.1 Logical Transformation

We commence with an intuitive, graph-agnostic transformation before introducing a graph-aware technique grounded on the concept of decomposition tree, which is the key to the optimization of graph pattern matching in the literature [26, 43].

Before proceeding, we introduce the concept of pattern decomposition that decomposes \mathcal{P} into two overlapping patterns, \mathcal{P}_1 and \mathcal{P}_2 , with shared vertices $V_o = V_{\mathcal{P}_1} \cap V_{\mathcal{P}_2}$ and shared edges

Table 1: Frequently used notations.

Notation	Definition
R	a relation or relational table
τ and $\tau.attr$	a tuple in a relation, and the value of an attribute of τ
$G(V, E)$	a property graph with V and E
$\mathcal{P}(V, E)$	a pattern graph with V and E
$id(\epsilon), \ell(\epsilon), \epsilon.attr$	the identifier, label, and the value of given attribute of a graph element ϵ
$N(u)$ and $N^E(u)$	neighbors and adjacent edges of u
GR	a graph relation
$M(GR, \mathcal{P}), M(\mathcal{P})$	matching \mathcal{P} on a graph relation GR or a graph G
$\pi_A, \sigma_\Psi, \bowtie$	projection, selection, and join operators over relations
$\hat{\pi}_{A*}, \hat{\bowtie}$	projection and join operators over graph relations
$\lambda_\ell^s(e), \lambda_\ell^t(e)$	the total functions for mapping tuples in an edge relation to source and target vertex relations

$E_o = E_{\mathcal{P}_1} \cap E_{\mathcal{P}_2}$. Denote $\mathcal{P} = \mathcal{P}_1 \cup \mathcal{P}_2$. Under the homomorphism semantics, the matching of \mathcal{P} can be represented as:

$$M(\mathcal{P}) = M(\mathcal{P}_1) \hat{\bowtie}_{V_o, E_o} M(\mathcal{P}_2), \quad (2)$$

where $\hat{\bowtie}$ is a natural join operator for joining two graph relations based on the common vertices and edges.

It is important to note that Eq. 2 is also applicable to alternative semantics, including isomorphism and non-repeated-edge [5]. To support these semantics, a special all-distinct operator can be applied as a filter to remove results that contain duplicate vertices and/or edges. The adoption of the all-distinct operator is compatible with all techniques in this paper.

3.1.1 Graph-agnostic Transformation. If the matching operator can be transformed into purely relational operations, the SPJM query becomes a standard SPJ query, which can then be optimized using existing relational optimizers (Sec. 4.1). This graph-agnostic approach is intuitive and easy to implement on top of existing relational databases, making it a straightforward choice in prototyped systems [1, 40, 41]. However, there is no theoretical guarantee that such a transformation is lossless in the context of RGMapping (Def. 1).

In this subsection, we bridge this gap by demonstrating the lossless transformation of the matching operator under RGMapping.

Consider a pattern graph \mathcal{P} and one of its edges $e = (u_s, u_t)$. According to the definition of the matching operator (Sec. 2.2), the graph edges and vertices that can be matched with e must have the labels $\ell(e)$, $\ell(u_s)$, and $\ell(u_t)$. We further denote the relations corresponding to these edges and vertices via RGMapping as $R_{\ell(e)}$, $R_{\ell(u_s)}$, and $R_{\ell(u_t)}$, respectively. Moreover, there must be total functions $\lambda_{\ell(e)}^s$ and $\lambda_{\ell(e)}^t$ for mapping tuples from $R_{\ell(e)}$ to $R_{\ell(u_s)}$ and $R_{\ell(u_t)}$, respectively. We define the following EVJoin relational operation regarding $\lambda_{\ell(e)}^s$ as:

$$R_{\ell(e)} \bowtie_{\ell(e)} R_{\ell(u_s)} = \{(\tau_e, \tau_s) \mid \tau_e \in R_{\ell(e)} \wedge \tau_s \in R_{\ell(u_s)} \wedge \lambda_{\ell(e)}^s(\tau_e) = \tau_s\}. \quad (3)$$

The EVJoin regarding $\lambda_{\ell(e)}^t$ is defined analogously. Although called EVJoin, the operation is associative like any relation join, meaning that the order in which the edge and vertex relations are joined does not affect the final result.

We have the following lemma¹.

LEMMA 1. *Under RGMapping, the matching operation in an SPJM query can be losslessly transformed into a sequence of relational joins involving n vertex relations and m edge relations.*

EXAMPLE 4. *Given pattern graph \mathcal{P} shown in Fig. 1(b), the matching operation $\mathcal{M}(\mathcal{P})$ can be converted to a sequence of join operations as follows. Without loss of generality, we start from \mathcal{P}_0 containing only the vertex u_{p_1} , and we have $R_0 = R_{\text{Person}}^1$ (note that the superscript 1 is used to differentiate relations of the same name). Next, we sequentially add the edges $e_1 = (u_{p_1}, u_{p_2})$, $e_2 = (u_{p_1}, u_m)$, and $e_3 = (u_{p_2}, u_m)$ to \mathcal{P}_0 , resulting in the following relations:*

$$\begin{aligned} R'_1 &= R_{\text{Person}}^1 \bowtie_{\text{person_id}=\text{pid1}} R_{\text{Knows}} \bowtie_{\text{pid2}=\text{person_id}} R_{\text{Person}}^2, \\ R'_2 &= R_{\text{Person}}^1 \bowtie_{\text{person_id}=\text{pid}} R_{\text{Likes}}^1 \bowtie_{\text{mid}=\text{message_id}} R_{\text{Message}}, \\ R'_3 &= R_{\text{Person}}^2 \bowtie_{\text{person_id}=\text{pid}} R_{\text{Likes}}^2 \bowtie_{\text{mid}=\text{message_id}} R_{\text{Message}}. \end{aligned}$$

Finally, we have $\hat{\pi}_{A*}(\mathcal{M}(\mathcal{P})) = R'_3 \bowtie R'_2 \bowtie R'_1 \bowtie R_0$. Note that R_{Person}^1 in R'_2 , as well as R_{Person}^2 and R_{Message} in R'_3 , are redundant and can be safely removed from the final join. By eliminating these redundant relations, we obtain a sequence of joins involving 3 vertex relations and 3 edge relations.

3.1.2 Graph-aware Transformation. We introduce a graph-aware transformation that incorporates key ideas from the literature on graph optimization. Following Eq. 2, we can recursively decompose \mathcal{P} , forming a tree structure called the *decomposition tree*. The tree has a root node that represents \mathcal{P} , and each non-leaf *intermediate* node is a sub-pattern (a subgraph of the pattern) $\mathcal{P}' \subset \mathcal{P}$, which has a left and right child node, denoted as \mathcal{P}'_l and \mathcal{P}'_r , respectively. The leaf nodes of the tree are called *Minimum Matching Components* (MMC), correspond to indivisible patterns directly solvable with specific physical operations as will be introduced in Sec. 3.2. The decomposition tree naturally forms a logical plan for solving $\mathcal{M}(\mathcal{P})$, as demonstrated in Fig. 2. For any non-leaf node \mathcal{P}' , there exists a relationship $\mathcal{M}(\mathcal{P}') = \mathcal{M}(\mathcal{P}'_l) \bowtie \mathcal{M}(\mathcal{P}'_r)$ according to Eq. 2. The plan allows for the recursive computation of the entire pattern.

¹Omitted proofs can be found in the full paper [6].

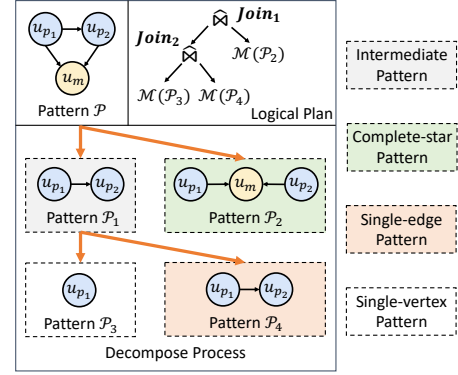


Figure 2: Example of decomposition trees and the corresponding logical plans.

Following state-of-the-art graph optimizers [26, 43], to guarantee a *worst-case optimal* execution plan [33], all intermediate sub-patterns in the decomposition tree must be induced subgraphs of \mathcal{P} . Furthermore, MMC is restricted to be a single-vertex pattern and a *complete star*. A star-shaped pattern is denoted as $\mathcal{P}(u; V_s)$, where u is the root vertex and V_s is the set of leaf vertices². In the decomposition tree, given $\mathcal{P}' = \mathcal{P}'' \cup \mathcal{P}(u; V_s)$, $\mathcal{P}(u; V_s)$ is a complete star if and only if it is a right child and $V_s \subseteq V_{\mathcal{P}'}$, meaning that the leaf vertices of the complete star must all be common vertices for the decomposition. A single-edge pattern is a special case of a complete star. The complete star logically represents the physical operations of EXPAND_INTERSECT, which will be discussed in Sec. 3.2. As shown in Fig. 2, a single-edge pattern, such as \mathcal{P}_3 , is further decomposed into a single-vertex pattern and the pattern itself, allowing the optimizer to select from which vertex the edge can be expanded.

REMARK 1. *The graph-aware transformation is fundamentally different from its graph-agnostic counterpart. While the graph-agnostic approach consistently converts pattern matching operations into relational joins between vertex and edge relations, the graph-aware transformation does not, due to the constraints imposed by pattern decomposition. While the graph-agnostic approach is straightforward, it has the following drawbacks:*

- **Graph-unaware Join Order:** It may cause the relational optimizer to change the order of joining vertex and edge relations, potentially missing opportunities to leverage graph indexes for the efficient computation of adjacent edges and vertices, as will be further discussed in Sec. 3.2.1.
- **Suboptimal Join Plans:** It generates plans that consistently reflect edge-based join plans that have been shown to be suboptimal in terms of worst-case performance [24].
- **Increased Search Space:** Compared to the graph-aware transformation, it can lead to an exponentially larger search space when computing optimal plans, which will be elaborated upon in the following subsection.

3.1.3 The Search Space: Graph-agnostic vs Graph-aware. After applying graph-agnostic transformations to the matching operator, the optimizer searches for the optimal join order. In contrast, applying graph-aware transformations leads to a search for the optimal

²Edge directions between u and V_s are not important, and we assume they all point from u to V_s .

decomposition tree. The search space for the graph-agnostic approach is clearly larger than that of the graph-aware approach, given the constraints imposed on the decomposition tree in the latter approach. However, the precise difference in search space complexity between the two approaches has not been rigorously analyzed. In this subsection, we derive a lower bound on the search space of the graph-agnostic approach and an upper bound on the search space of the graph-aware approach to demonstrate that the latter is exponentially more efficient in terms of search space.

According to Lemma 1, the graph-agnostic approach must produce a sequence of joins among n vertex relations and m edge relations, which is equivalent to searching for the optimal join order among $n + m$ relations. The following lemma establishes a lower bound on the search space complexity for this problem:

LEMMA 2. *Given a join among $n + m$ relations, the search space for determining the optimal join order is at least $\Omega(4^{m+n-1})$.*

In contrast, when the graph-aware transformation is applied, the search space is equivalent to the number of possible decomposition trees. Despite the numerous works proposed to optimize graph pattern matching [26, 31, 43], to the best of our knowledge, the search space of this optimization problem has not been thoroughly analyzed. The following lemma provides an upper bound on the search space complexity for the graph-aware approach:

LEMMA 3. *The search space for determining the optimal decomposition tree for pattern \mathcal{P} is at most $O(4^{n-1})$.*

We ultimately obtain the following theorem.

THEOREM 1. *The graph-aware transformation yields a search space that is exponentially smaller than that of the graph-agnostic transformation, for optimizing the matching operator in an SPJM query.*

PROOF. According to Lemma 2 and Lemma 3, we have $\frac{\Omega(4^{m+n-1})}{O(4^{n-1})} = \Omega(4^m)$. \square

REMARK 2. *It is worth noting that the search space after applying the graph-aware transformation, as given in Lemma 3, does not depend on m , the number of edges in the pattern. This is mainly due to the constraints imposed on the decomposition tree, which require each intermediate pattern to be an induced sub-pattern. Consequently, given any subset of vertices, enumerating edges among these vertices is not necessary. To further verify the findings in Theorem 1, we conducted a micro-benchmark experiment using a path graph with m edges. We programmed an enumerator to explore the search space of both the graph-agnostic and graph-aware approaches while varying m . The results, shown in Fig. 3, confirm the significant difference in search space size between the two approaches.*

3.2 Physical Implementation

In the graph view, given a vertex v , it is efficient to obtain its adjacent edges and vertices (i.e., neighbors). However, in the relational view, such adjacency relationships between vertices and edges are not directly stored in relations but must be computed via the EVJoin operations (Eq. 3). While there are multiple ways to construct the graph view in the literature [18, 39], we refer to the method introduced in GrainDB [21], which is free from materializing the graph. This approach avoids the extra storage cost associated with





Figure 3: Search Space Comparison.

likes_id	pid	mid	date	pid_rowid	mid_rowid
l_1	p_1	m_1	2024-03-31	0	0
l_2	p_2	m_1	2024-03-28	1	0
l_3	p_2	m_2	2024-03-20	1	1
l_4	p_3	m_2	2024-03-21	2	1

(a) The EV-Index on R_{Likes}



(b) The VE-Index on R_{Person} in the CSR Format

Figure 4: The graph index constructed among relations R_{Person} , R_{Likes} and $R_{Message}$ in Fig. 1(a).

graph materialization and ensures compatibility with the relational context. Specifically, GrainDB introduces an indexing technique called pre-defined join to improve the performance of join operations. As the pre-defined join essentially materializes the adjacency relationships, we treat it as a *graph index* in this work.

3.2.1 *Graph Index.* As shown in Fig. 4, given the three relations R_{Person} , R_{Likes} , and $R_{Message}$, the complete information of “Person likes messages” can be obtained by conducting the join:

$$R_{Person} \bowtie_{\text{person_id} = \text{pid}} R_{Likes} \bowtie_{\text{mid} = \text{message_id}} R_{Message}.$$

GrainDB introduces two kinds of indexes to the relational tables to efficiently process the join: the EV-index and the VE-index. The EV-index, shown in Fig. 4(a), is constructed by appending extra columns to the table R_{Likes} . The column “pid_rowid” stores the row ID of the corresponding tuple in the table R_{Person} , denoted as $\text{rid}(\tau_p)$, where $\tau_p \in R_{Person}$. Similarly, the column “mid_rowid” stores the row ID of the corresponding tuple in the table $R_{Message}$, denoted as $\text{rid}(\tau_m)$, where $\tau_m \in R_{Message}$. These row ids help quickly route a tuple $\tau_l \in R_{Likes}$ to the joinable tuples τ_p and τ_m without additional operations like hash-table lookup or sorting.

The VE-index in Fig. 4(b) is created on R_{Person} for efficiently computing its “liked messages”. For each tuple $\tau_p \in R_{Person}$, the VE-index records the row ids of tuples $\tau_l \in R_{Likes}$ and the corresponding $\tau_m \in R_{Message}$ that are joinable with τ_p . In the graph view, treating “Person-[Likes]->Messages” as an edge of a property graph, the VE-index maintains the adjacent edges and vertices of each person.

We can adopt GrainDB’s approach to construct the graph indexes during the RGMMapping process. Given an edge relation R_e and its associated vertex relations R_{v_s} and R_{v_t} , the EV-index can be constructed on R_e for each tuple $\tau_e \in R_e$ by including $\text{rid}(\lambda_e^s(\tau_e))$ and $\text{rid}(\lambda_e^t(\tau_e))$, which are the row ids of the corresponding tuples

in R_{v_s} and R_{v_t} , respectively. Meanwhile, the VE-index can be constructed on R_{v_s} for each tuple $\tau_{v_s} \in R_{v_s}$ by including the row ids of all tuples $\tau_e \in R_e$ such that $\lambda_e^s(\tau_e) = \tau_{v_s}$, along with the row ids of the corresponding tuples $\tau_{v_t} \in R_{v_t}$ such that $\lambda_e^t(\tau_e) = \tau_{v_t}$. The construction of VE-index on R_{v_t} is analogous.

3.2.2 The Graph-Aware Execution Plan. We delve into the physical implementation of the execution plan provided by the graph-aware method for solving $\mathcal{M}(\mathcal{P})$. The entry point of the plan is always matching a single-vertex pattern \mathcal{P}_u , which is one of the leaf nodes in the decomposition tree.

The implementation of $\mathcal{M}(\mathcal{P}_u)$ is straightforward: scanning the corresponding vertex relation $R_{\ell(u)}$ and encoding each tuple as a graph vertex object that contains its ID and label (mandatory) and necessary attributes. The row ID of the tuple in the relation can be directly used as the ID. To ensure globally uniqueness, the name of the relation can be incorporated as a prefix of the ID. Advanced encoding techniques are necessary for production use, but they are beyond the scope of this paper.

The plan is then constructed in a bottom-up manner. As shown in Fig. 2, there are three fundamental cases to consider when implementing the plan.

Case I: Solving $\mathcal{M}(\mathcal{P}') = \mathcal{M}(\mathcal{P}'_l) \bowtie_{V_o, E_o} \mathcal{M}(\mathcal{P}'_r)$, where \mathcal{P}'_l and \mathcal{P}'_r are both intermediate patterns in the decomposition tree. The implementation of such a join is similar to a conventional relational join. The join is constrained to a natural join, where the join condition is simply the equality of the common vertices V_o and edges E_o between \mathcal{P}'_l and \mathcal{P}'_r . During the implementation of the join, the identifiers of the vertices and edges can serve as the keys for comparison. Note that the input and output of the join are both graph relations, which will not be projected into relational tuples until the last stage that obtains the results $\mathcal{M}(\mathcal{P})$.

Case II: Solving $\mathcal{M}(\mathcal{P}') = \mathcal{M}(\mathcal{P}'_l) \bowtie_{u_s} \mathcal{M}(\mathcal{P}_e)$, where \mathcal{P}_e is a single-edge pattern, and u_s is the source vertex in \mathcal{P}'_l from which the edge $e = (u_s, u_t)$ is expanded. Note that it's not possible for both u_s and u_t to be in \mathcal{P}'_l , as it would violate the fact that \mathcal{P}'_l is either a single vertex or an induced sub-pattern.

When there is no graph index, $\mathcal{M}(\mathcal{P}_e)$ is computed via $R_{\ell(u_s)} \bowtie_{e \in V} R_{\ell(e)} \bowtie_{e \in V} R_{\ell(u_t)}$. This case is then reduced to Case I.

When graph indexes exist, the implementation is handled by the physical operators of EXPAND_EDGE and GET_VERTEX. For each tuple $\tau \in \mathcal{M}(\mathcal{P}'_l)$, $\tau.u_s$ must record a graph vertex v_s that matches u_s in the pattern \mathcal{P}'_l . The EXPAND_EDGE operator looks up the VE-index of v_s , which allows it to efficiently compute v_s 's adjacent edges (more precisely, it's the corresponding edge tuples). Furthermore, the GET_VERTEX operator is used to obtain the matched vertex v_t that is connected to v_s via the previous matched edges, which can be achieved by looking up the EV-index of the matched edges. By combining the results of EXPAND_EDGE and GET_VERTEX, the tuple of $(\tau, \mathcal{N}^E(v_s), \mathcal{N}(v_s))$ is rendered. For example, in Fig. 4(b), if we apply EXPAND_EDGE and GET_VERTEX to a tuple τ from v_{p_2} , the result $(\tau, [e_{l_2}, e_{l_3}], [v_{m_1}, v_{m_2}])$ is returned. Furthermore, to obtain $\mathcal{M}(\mathcal{P}')$, we flatten the adjacent edges and vertices and pair them up. In the case of $(\tau, [e_{l_2}, e_{l_3}], [v_{m_1}, v_{m_2}])$, two tuples (τ, e_{l_2}, v_{m_1}) and (τ, e_{l_3}, v_{m_2}) are generated.

In real-life scenarios, a vertex may be adjacent to multiple types of edges. For example, in Fig. 1, a Person vertex can be connected to

both Likes and Knows edges. To handle such cases, we can record edge's ID instead of just the row ID of the tuple. Given that the edge's ID is a combination of its label and the tuple's row ID, the adjacent edges of a specific label can be easily obtained from the VE-Index.

Case III: Solving $\mathcal{M}(\mathcal{P}') = \mathcal{M}(\mathcal{P}'_l) \bowtie_{V_s, E_s} \mathcal{M}(\mathcal{P}(u; V_s))$, where pattern $\mathcal{P}(u; V_s)$ is a complete k -star with $V_s = \{u_1, \dots, u_k\}$.

When there is no graph index, solving Case III involves continuously joining $|V_s|$ single-edge patterns.

When graph indexes are available, the EXPAND_INTERSECT operator can be used to efficiently compute the join. Given a tuple $\tau \in \mathcal{M}(\mathcal{P}'_l)$, let $\{v_1, \dots, v_k\}$ be the vertices in τ that match the leaf vertices $\{u_1, \dots, u_k\}$ in the complete star $\mathcal{P}(u; V_s)$. The vertices that can match the root vertex u of the star must be the common neighbors of all the leaf vertices.

Consequently, for the tuple τ , the physical EXPAND_INTERSECT operator performs the following steps:

- (1) For each leaf vertex $u_i \in V_s$ ($1 \leq i \leq k$), apply the EXPAND_EDGE and GET_VERTEX operators to obtain the adjacent edges and neighbors of the corresponding vertices v_i respectively.
- (2) Compute the intersections of all adjacent edges and neighbors returned by the EXPAND_EDGE and GET_VERTEX operators.
- (3) Return a new tuple as follows; for the sake of simplicity, the details of the edges are omitted: $(\tau, \bigcap_{1 \leq i \leq k} \mathcal{N}(v_i))$.

Note that the above step (1) and (2) can be computed in a pipeline manner, following a certain order of among the leaf vertices. Similar to Case II, we flatten the common edges and vertices and pair them up to obtain the final result.

EXAMPLE 5. Given \mathcal{P} in Fig. 2, a decomposition tree and its corresponding logical plan are presented. We illustrate the physical implementation of $\mathcal{M}(\mathcal{P}_1) \bowtie \mathcal{M}(\mathcal{P}_2)$ using EXPAND_INTERSECT when a graph index is available. Consider the tuple $(v_{p_1}, e_{k_1}, v_{p_2})$ from $\mathcal{M}(\mathcal{P}_1)$ as an example. First, the EXPAND_EDGE and GET_VERTEX operators are applied to obtain the adjacent edges and neighbors of v_{p_1} and v_{p_2} , resulting in

$$(v_{p_1}, e_{k_1}, v_{p_2}, [e_{l_1}], [v_{m_1}]), \text{ and} \\ (v_{p_1}, e_{k_1}, v_{p_2}, [e_{l_2}, e_{l_3}], [v_{m_1}, v_{m_2}]).$$

Next, the intersection process is conducted. Since $\mathcal{N}(v_{p_1}) \cap \mathcal{N}(v_{p_2}) = [v_{m_1}]$, the edges in both sets that have v_{m_1} as the target vertex are retained, resulting in $(v_{p_1}, e_{k_1}, v_{p_2}, [(e_{l_1}, e_{l_2}, v_{m_1})])$. Finally, the tuple is flattened to $(v_{p_1}, e_{k_1}, v_{p_2}, e_{l_1}, e_{l_2}, v_{m_1})$.

4 THE CONVERGED OPTIMIZATION FRAMEWORK

This section presents RelGo, a converged relational/graph optimization framework designed to optimize the query processing of SPJM queries. We begin by introducing a naive solution built upon the graph-agnostic method for solving the matching operator. We then delve into the converged workflow of RelGo, which leverages the graph-aware method for solving the matching operator and introduces a complete workflow that aims to integrate techniques from both relational and graph optimization modules.

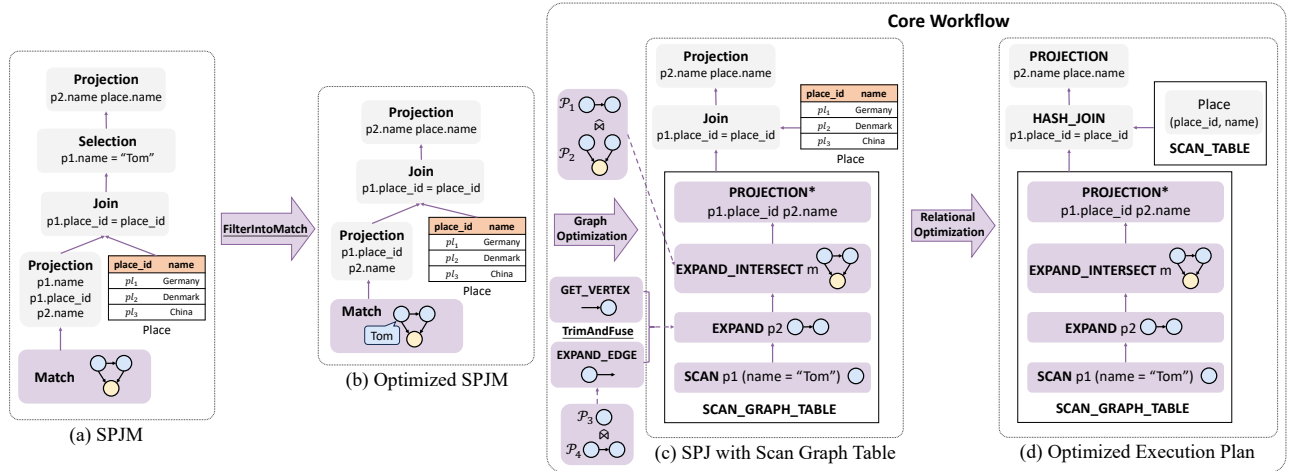


Figure 5: The converged optimization workflow

4.1 Graph-Agnostic Approach

The graph-agnostic approach is straightforward: it applies the graph-agnostic transformation for the matching operator in an SPJM query into a series of relational operations (Lemma 1), effectively converting the SPJM query into an SPJ query. The resulting SPJ query can then be optimized by any existing relational optimizer, producing an execution plan. As an improvement, if a graph index (Sec. 3.2.1) is available, certain hash-join operators in the execution plan can be replaced by the predefined-join operator, as discussed in GrainDB [21]. The main advantage of this solution is its easy integration with any existing relational database. However, it suffers from two significant drawbacks discussed in Remark 1.

4.2 The Converged Approach

As illustrated in Fig. 5, the core workflow of the RelGo framework consists of two components: *graph optimization* and *relational optimization*. The graph optimization is responsible for handling the graph component in an SPJM query, leveraging graph optimization techniques to determine the optimal decomposition tree of the matching operator. On the other hand, the relational optimization takes over to optimize the relational component in the query. The order in which these two components are applied is not strictly defined. However, for the purpose of our discussion, we will first focus on the graph optimization and then proceed to the relational optimization. In addition to the core workflow, we further explore heuristic rules that highlight the non-trivial interplay between the relational and graph components in an SPJM query.

4.2.1 The Graph Optimization. In this work, we adopt the graph optimization techniques developed in GLogS [26]. However, it is crucial to note that GLogS was originally designed for native graph data, whereas our framework deals with relational data, which necessitates a careful adaptation of GLogS's techniques to the relational setting.

GLogue Construction. GLogS is built upon a data structure called GLogue, which is essentially a graph $\mathbb{G}_{\mathcal{P}}(V, E)$. In this graph, each vertex represents a pattern \mathcal{P}' consisting of up to k vertices (typically, $k = 3$) that has non-empty matched instances in the original

graph. There is an edge from \mathcal{P}'' to \mathcal{P}' , if there is a decomposition tree where \mathcal{P}'' is a child node of \mathcal{P}' .

Each vertex \mathcal{P}' in the GLogue maintains $|\mathcal{M}(\mathcal{P}')|$, denoting the cardinality of the pattern. To reduce computation costs, GLogS employs a sparsification technique that constructs a smaller sub-graph G' . The cardinality of the pattern can then be estimated using $|\mathcal{M}_{G'}(\mathcal{P}')|$, computed based on the subgraph G' . In our work, we adapt this sparsification technique to construct the GLogue. To do so, we sample a subset of vertex and edge relations during the RGMMapping process. Once the subset of relations is obtained, they can serve as the input tables to the techniques presented in [39] for constructing the sparsified graph G' .

Cost Calculation. The optimization process is essentially searching for the execution plan that incurs the minimal cost. Let the cost of an execution plan Φ for computing $\mathcal{M}(\mathcal{P})$ be $\text{Cost}_{\Phi}(\mathcal{P})$.

Consider $\mathcal{M}(\mathcal{P}') = \mathcal{M}(\mathcal{P}'_l) \bowtie \mathcal{M}(\mathcal{P}'_r)$ as an intermediate computation in an execution plan. We have:

$$\text{Cost}_{\Phi}(\mathcal{P}') = \text{Cost}_{\Phi_l}(\mathcal{P}'_l) + \text{Cost}_{\Phi_r}(\mathcal{P}'_r) + \text{Cost}(\bowtie),$$

where Φ_l and Φ_r are the execution plans for computing $\mathcal{M}(\mathcal{P}'_l)$ and $\mathcal{M}(\mathcal{P}'_r)$, respectively, and $\text{Cost}(\bowtie)$ is the cost of the join operation.

When a graph index is available, there are three physical implementations of \bowtie , depending on the type of \mathcal{P}'_r , and the calculation of $\text{Cost}(\bowtie)$ differs accordingly:

- If \mathcal{P}'_r is a single-edge pattern, \bowtie is implemented using the **EXPAND_EDGE** operator followed by **GET_VERTEX**. The cost is calculated based on the cardinality of $\mathcal{M}(\mathcal{P}'_l)$ (can be looked up in the GLogue) and the average degree of the graph, namely $|\mathcal{M}(\mathcal{P}'_l)| \times \bar{d}$.
- If \mathcal{P}'_r is a complete star pattern, \bowtie is implemented using the **EXPAND_INTERSECT** operator. The cost is calculated based on the cardinality of $\mathcal{M}(\mathcal{P}'_l)$ and the average intersection size of the neighbors of the vertices being intersected, which is maintained on the corresponding edge from \mathcal{P}' to \mathcal{P}'_l in GLogue.
- If \mathcal{P}'_r is any arbitrary pattern, \bowtie is implemented as a **HASH_JOIN**. The cost is calculated as the product of the cardinalities of the two relations being joined, i.e., $\text{Cost}(\bowtie) = |\mathcal{M}(\mathcal{P}'_l)| \times |\mathcal{M}(\mathcal{P}'_r)|$.

In the absence of a graph index, HASH_JOIN is used for the entire plan, where the cost is simply the product of the cardinalities of the two relations being joined.

Plan Computation. Searching for the optimal execution plan in RelGo remains the same as in GLogS. The optimal plan is obtained by searching for the shortest path in the GLogue from the single-vertex pattern to the queried pattern. Fig. 5(c) demonstrates a physical plan for matching the given triangle pattern when a graph index is present. The plan reflects the example in Example 5, with one exception: the pair of EXPAND_EDGE and GET_VERTEX operators is fused into a single EXPAND operator, which will be discussed as a heuristic rule called TrimAndFuseRule.

4.2.2 The Relational Optimization. Once the graph optimizer has computed the optimal execution plan for $\mathcal{M}(\mathcal{P})$, the next step is to integrate this plan with the remaining relational operators in the SPJM query. The relational optimization is responsible for optimizing these remaining operators, which are all relational operators. Any existing relational optimizer can be used for this purpose.

To prevent the relational optimizer delve into the internal details of the graph pattern matching process, we introduce a new physical operator called SCAN_GRAPH_TABLE, as shown in Fig. 5(c), which encapsulates the $\hat{\pi}_{A^*}$ operator and the optimal execution plan for $\mathcal{M}(\mathcal{P})$. The SCAN_GRAPH_TABLE operator acts as a bridge between the graph and relational components of the query. From the perspective of the relational optimizer, SCAN_GRAPH_TABLE behaves like a standard SCAN operator, providing a relational interface to the matched results.

4.2.3 Heuristic Optimization Rules. In real-life use cases, heuristic rules may involve non-trivial interactions between the relational and graph components of an SPJM query. We explore two representative rules, FilterIntoMatchRule and TrimAndFuseRule, which can be applied at different stages of the optimization process to improve query performance.

FilterIntoMatchRule. To elaborate the rule, we extend the definition of a pattern (\mathcal{P}, Ψ) , introducing constraints within Ψ . For example, constraints can specify predicate d such as $\text{id}(v_1) = p_1$ for a vertex v_1 , or $e_1.\text{date} > "2024-03-31"$ for an edge e_1 . With the constraints defined, any matching result of \mathcal{P} must have the corresponding vertices and edges adhering to the predicates.

Often while writing queries, users may not specify constraints on the pattern but rather use the selection operator after the matching results have been projected into the relational relation, which can be described as:

$$\sigma_{d'_{v_a}}(\hat{\pi}_{v_a \rightarrow v_a, \dots} \mathcal{M}(\mathcal{P}))$$

The predicate d'_{v_a} defines a predicate in terms of an attribute of the pattern vertex that is projected by $\hat{\pi}$ from the matched results. The motivation example in Example 1 illustrates such a case, where the selection predicate $\text{g.p1_name} = \text{"Tom"}$ is applied to the pattern vertex v_{p_1} . There is wasteful computation if the selection is applied after the costly pattern matching. A more efficient approach is to push the selection predicate down into the matching operator. The FilterIntoMatchRule is formally defined as:

$$\sigma_{\Psi'}(\hat{\pi}_{v_a \rightarrow v_a, \dots} \mathcal{M}(\mathcal{P})) \equiv \sigma_{\Psi'}(\hat{\pi}_{v_a \rightarrow v_a, \dots} \mathcal{M}((\mathcal{P}, \{d_v\}))),$$

where $\Psi' = \Psi \setminus \{d'_{v_a}\}$, and $\{d_v\}$ is the corresponding constraints that are appended to the pattern \mathcal{P} .

It is recommended to apply the FilterIntoMatchRule before graph optimization, as this allows the optimizer to leverage the pushed-down constraints to recalculate the cost, potentially generating more efficient execution plans. Fig. 5(b) showcases the effects of applying the FilterIntoMatchRule, where the selection predicate $\text{g.p1_name} = \text{"Tom"}$ is pushed down into the matching operator.

TrimAndFuseRule. The TrimAndFuseRule aims to streamline a query plan by merging the EXPAND_EDGE and GET_VERTEX operators which are commonly coupled in the implementation of matching operations, into a single EXPAND operator that retrieves the neighboring vertices directly. However, such a fusion is permissible solely when the output edges by EXPAND_EDGE are deemed unnecessary, so this rule further incorporates a preceded field trim step. Specifically, the field trimmer would examine whether any subsequent relational processes rely on these edges, such as utilizing them for property projections or for filtering based on their attributes. If no such operations are found, the edges can be trimmed. Furthermore, the field trimmer would also consider a special case that the edges might be projected in the SCAN_GRAPH_TABLE operator as part of the matching results, but are subsequently unused in relational operations. In such cases, the edges can be trimmed as well. After the field trim step, if the output edges are trimmed, the EXPAND_EDGE operator can be fused with the GET_VERTEX operator to form a single EXPAND operator, which can directly retrieve the neighboring vertices efficiently by looking up the VE-index of the source vertex when the graph index is available.

4.3 System Implementation

We engineered the frontend of RelGo in Java and built it upon Apache Calcite [14], a widely used open-source framework for data management, to utilize its robust relational query optimization infrastructure. Firstly, we enhanced Calcite's SQL parser to recognize SQL/PGQ extensions, specifically to parse the GRAPH_TABLE clause. We created a new ScanGraphTableRelNode that inherits from Calcite's core RelNode class, translating the GRAPH_TABLE clause into this newly defined operator within the logical plan. Following the formation of the logical plan, the frontend invokes the converged optimizer to generate the optimal physical plan. For the relational-graph interplay optimizations, we incorporate heuristic rules such as FilterIntoMatchRule and TrimAndFuseRule into Calcite's rule-based HepPlanner, by specifying the activation conditions and consequent transformations of each rule. For more nuanced optimization, we rely on the VolcanoPlanner, the cost-based planner in Calcite, to optimize the ScanGraphTableRelNode. We devised a top-down search algorithm that assesses the most efficient physical plan based on a cost model outlined in Sec. 4.2.1, combined with high-order statistics from GLogue for more accurate cost estimation. For the remaining relational operators in the query, we leverage Calcite's built-in optimizer, which already includes comprehensive relational optimization techniques. Lastly, the converged optimizer outputs an optimized and platform-independent plan formatted with Google Protocol Buffers (protobuf) [36]. This ensures the adaptability of the RelGo framework's output to a variety of backend database systems.

We have developed the backend of the RelGo framework using C++, on top of DuckDB as the underlying relational execution engine, to demonstrate the optimization capabilities that RelGo

can bring. We integrated graph index support in GrainDB [21]. With graph index, the EXPAND, EXPAND_EDGE and GET_VERTEX operators can be optimized by directly using the predefined join in GrainDB. Note that we craft a new join on DuckDB called *EI-Join* for the support of EXPAND_INTERSECT. Without graph index, the HASH_JOIN operator is used throughout the entire plan. To execute the optimized plans within DuckDB, we introduced a runtime module that translates the optimized physical plan into a sequence of DuckDB/GrainDB-compatible executable operators. This runtime module essentially bridges the gap between the optimized plans produced by the RelGo framework and DuckDB’s execution engine, thereby validating RelGo’s practicality and its potential to improve the performance of SPJM queries on an established relational database system.

Dataset	V	E	Disk Space Usage
G_{sf10}	29,987,835	88,317,856	8.9G
G_{sf30}	88,789,833	278,652,443	28.0G
IMDB	14,431,946	59,758,241	3.7G

Table 2: Statistics of the datasets. In detail, $|\cdot|$ represents the number of elements in \cdot .

5 EVALUATION

5.1 Experimental Settings

Benchmarks. Our experiments leverage two widely used benchmarks to assess system performance, as follows:

- **LDBC SNB.** We follow the Linked Data Benchmark Council (LDBC) social network benchmark [27] to test the system. We use two commonly-adopted datasets with scale factors of 10 and 30, denoted as G_{sf10} and G_{sf30} respectively, generated by the official LDBC Data Generator. We select 10 queries from the LDBC Interactive workload for evaluation, denoted as $IC[1, \dots, 9, 11, 12]$, with 10, 13, and 14 excluded since they involve either pre-computation or shortest-path that are not supported. To accommodate queries containing variable-length paths [21], we followed [21] to slightly modify them by separating each query into multiple individual queries with fixed-length paths. Each of these modified queries is denoted with a suffix “-l”, where l represents the length of the fixed-length path. In addition, we carefully designed two sets of queries for the comprehensiveness of evaluation, including (1) $QR[1 \dots 4]$ to test the effectiveness of heuristic rules FilterIntoMatchRule and TrimAndFuseRule in RelGo, and (2) $QC[1 \dots 3]$, comprising three typical patterns with cycles including triangle, square, and 4-clique, to assess the efficiency of EXPAND_INTERSECT introduced in Sec. 3.2.
- **JOB.** The Join Order Benchmark (JOB) [28] on Internet Movie Database (IMDB) is adopted. We select the variants marked with “a” of all JOB queries, referred to as $JOB[1 \dots 33]$, without loss of generality. These queries are primarily designed to test join order optimization, with each query containing an average of 8 joins.

The details of the dataset statistics are summarized in Table 2. We manually implement the queries using SQL/PGQ, which are presented in the full paper [6]. In addition, we adhered to the same procedures for data loading and graph index construction as in [21].

Compared Systems. To ensure a fair comparison, all the following systems use DuckDB as the underlying relational execution engine. They differ in the adoption of the optimizer.

- DuckDB [2]: This system optimizes queries using the graph-agnostic approach, leveraging DuckDB’s built-in optimizer as described in Sec. 4.1. It serves as the naive baseline for extending a relational database system to support SPJM.
- GrainDB [21]: This system uses same optimizer as DuckDB but employs the graph index (Sec. 3.2.1) for query execution. It acts as the baseline to demonstrate that solely using graph index is insufficient for optimizing SPJM.
- RelGo: This system optimizes queries using the converged optimizer presented in Sec. 4.2 and utilizes the graph index for query execution. It demonstrates the full range of techniques introduced in this paper. There are some variants of RelGo for verifying the effectiveness of the proposed techniques, which will be introduced in the corresponding experiments.

Configurations. Our experiments were conducted on a server equipped with an Intel Xeon E5-2682 CPU running at 2.50GHz and 251GB of RAM, with parallelism restricted to a single thread. For a comprehensive performance analysis, each query from the LDBC benchmark was run 50 times using the official parameters, while each query from the JOB benchmark was executed 10 times. We report the average time cost for each query to mitigate potential biases. We imposed a timeout limit of 10 minutes for each query, and queries that fail to finish within the limit are marked as OT.

5.2 Micro Benchmarks on RelGo

In this subsection, we conducted three micro benchmarks to evaluate the effectiveness of RelGo, including assessing the efficiency of the optimizer, testing its advanced optimization strategies, and examining its effectiveness in optimizing join order.

Optimization Efficiency Evaluation. First, we assessed the optimization efficiency by comparing RelGo with GrainDB[21]. We compared the optimization time of RelGo and GrainDB, and also evaluated the execution time for their optimized plans as a measure of the plan quality. We considered end-to-end time as optimization time plus execution time. We randomly selected two subsets of the LDBC and JOB queries, and conducted the experiments on G_{sf30} and IMDB datasets respectively.

The results, shown in Fig. 6, reveal that RelGo significantly outperforms GrainDB in terms end-to-end time. Specifically, RelGo achieves an average speedup of 6.5 \times over GrainDB on the G_{sf30} dataset and 2.2 \times on the IMDB dataset. However, it is worth noting that RelGo incurs a slightly higher optimization cost compared to GrainDB. Although RelGo theoretically has a narrower search space, as analyzed in Sec. 3.1.3, GrainDB benefits from DuckDB’s optimizer, which includes very aggressive pruning strategies.

Despite the slightly higher optimization cost, the quality of the optimized plans generated by RelGo is generally superior to those produced by GrainDB. This is evident in the results of execution time, where RelGo outperforms GrainDB by an average of 9.5 \times on the G_{sf30} dataset and 2.8 \times on the IMDB dataset.

For fair comparison, in the subsequent experiments, we evaluate the efficiency of different systems using the end-to-end time.

Advanced Optimization Strategies. In this experiment, we assessed the advanced optimization strategies in RelGo, including the




Figure 6: Experiments on optimization and execution cost




Figure 7: Efficiency comparison of RelGo and RelGoNoRule




Figure 8: Efficiency comparison of RelGo and RelGoNoEI

heuristic FilterIntoMatchRule and TrimAndFuseRule, and the optimized implementation of EXPAND_INTERSECT operator that aims to improve the efficiency of complete star join.

We began by testing FilterIntoMatchRule and TrimAndFuseRule, two representative heuristic rules employed in RelGo that capture optimization opportunities at the interplay of relational and graph optimizations. We conducted experiments on the G_{sf10} and G_{sf30} datasets, using QR[1] and QR[2] to evaluate FilterIntoMatchRule, and QR[3] and QR[4] to test TrimAndFuseRule. The results, depicted in Fig. 7, compared the performance of RelGo with and without applying these rules, denoted as RelGo and RelGoNoRule, respectively. The results demonstrate that FilterIntoMatchRule significantly improves query performance, providing an average speedup of 299.4 \times on G_{sf10} and 699.8 \times on G_{sf30} . After applying TrimAndFuseRule, query execution is accelerated by an average of 2.0 \times on G_{sf10} and 2.3 \times on G_{sf30} . These findings suggest that the heuristic rules, particularly FilterIntoMatchRule, are highly effective in enhancing query execution efficiency.

Next, we evaluated the effectiveness of the EXPAND_INTERSECT, which focuses on improving the efficiency of complete star join. Without this optimization strategy, the EXPAND_INTERSECT operator would be implemented as a traditional multiple join, and we denote this variant as RelGoNoEI. We conducted this experiment with queries QC[1...3] that contain cycles and compared the performance of RelGo and RelGoNoEI. The performance results, depicted in Fig. 8, suggest that, compared to RelGoNoEI, RelGo achieves an average speedup of 1.22 \times on G_{sf10} and 1.31 \times on G_{sf30} (excluding QC[3]). Notably, for QC[3], which is a complex 4-clique, the plans optimized by RelGoNoEI confront an out-of-memory (OOM) error. The results of this experiment indicate that

the EXPAND_INTERSECT operator with an optimized implementation not only enhances query performance but also significantly reduces the spatial overhead.




Figure 9: Experiments that compare the efficiency of RelGo and RelGoHash.

Efficiency of Join Order. In this experiment, we compared RelGo with GrainDB and DuckDB, focusing on the efficiency of the join order. For this purpose, we introduced a variant of RelGo called RelGoHash, which optimizes the plan in a converged manner like RelGo but deliberately bypasses the use of graph index. We randomly selected 10 queries from the JOB benchmark and showed the performance results in Fig. 9. The results demonstrate that RelGo outperforms GrainDB on all the queries, accelerating the execution time by factors ranging from 1.2 \times to 7.5 \times , with an average speedup of 3.31 \times . Additionally, the plans optimized with RelGoHash are at least as good as those optimized by DuckDB, achieving an average speedup of 1.4 \times . The effectiveness of RelGo and RelGoHash stems from their use of advanced graph-aware optimization techniques in optimizing the matching operator, resulting in good join order and thus robust performance regardless of the presence or absence of graph index. It is worth noting that RelGo does not always generate plans with the absolute best join orders, as it relies on the estimated cost of the plans. However, its optimized plans generally remain competitive in most cases, thanks to its integration of GLogue that use high-order statistics for cost estimation.

5.3 Comprehensive Experiments

We conducted comprehensive experiments on the LDBC and JOB benchmarks to comprehensively evaluate the performance of RelGo compared to GrainDB and DuckDB. The experimental results, presented in Fig. 10, demonstrate that execution plans optimized by RelGo consistently outperform those optimized by GrainDB or DuckDB on both benchmarks. Specifically, for the LDBC benchmark, the execution time of the plans optimized by RelGo is about 2.2 \times and 9.1 \times faster on average than those generated by GrainDB and DuckDB, respectively, on G_{sf10} , and about 4.10 \times and 14.72 \times faster on G_{sf30} . It is important to note that RelGo is especially effective for queries containing cycles, which can benefit more from graph optimizations. For example, in query IC[7], which contains a cycle, RelGo outperforms GrainDB and DuckDB by 20.1 \times and 42.9 \times , respectively, on G_{sf30} . Conversely, the JOB benchmark, established for assessing join optimizations in relational databases, lacks any cyclic-pattern queries. Despite this, RelGo still achieves better performance compared to GrainDB and DuckDB, with an average speedup of 2.1 \times and 5.1 \times , respectively.

The experiment results reflect our discussions in Sec. 3.1.2. We summarize the superiority of RelGo as follows. First, RelGo is designed to be aware of the existence of graph index in query optimization and can leverage the index to effectively retrieve adjacent edges and vertices. In contrast, for GrainDB, relational optimizers can occasionally alter the order of EVJoin operations, rendering

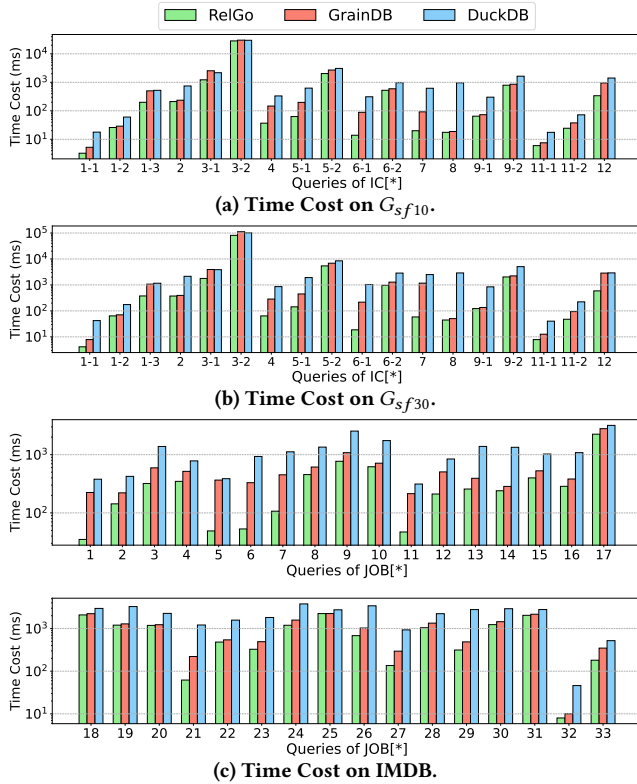


Figure 10: Results of the comprehensive experiments.

graph index ineffective. DuckDB, on the other hand, does not consider graph index in query optimization and executes queries using conventional hash joins, which are often less efficient compared to graph-aware approaches. Second, by incorporating a matching operator in SPJM queries to capture the graph query semantics, RelGo is able to leverage advanced graph optimization techniques to optimize the matching operator. These techniques include using high-order statistics to estimate the cost of plans more accurately and employing worst-case optimal join implementations to optimize cyclic patterns. In contrast, DuckDB and GrainDB cannot benefit from these graph-specific optimizations, which may lead to suboptimal plans and inefficient execution. Third, RelGo considers optimization opportunities across both graph and relational query semantics, introducing effective heuristic rules such as FilterIntoMatchRule and TrimAndFuseRule. These rules can significantly improve the efficiency of the generated plans.

6 RELATED WORK

Query Optimization for Relational Databases. Various studies of query optimization for relational databases were proposed to find the optimal join order. Ibaraki et al. [20] proved the NP-complexity of the join order optimization problem and designed an efficient algorithm with a time complexity of $O(n^2 \log n)$ to optimize tree queries. Then, Krishnamurthy et al. [23] optimized the algorithm and proposed an algorithm with a time complexity of $O(N^2)$ by reusing the computation results. Moreover, Haffner et al. [16] converted the problem of join order optimization to that of finding the shortest path on directed graphs and solved the problem with

the A* algorithm. Kossmann et al. [22] summarized the methods to optimize queries with data dependencies, such as uniqueness constraints, foreign key constraints, and inclusion dependencies. All these techniques can be orthogonally adopted in RelGo’s relational optimization.

Query Optimization for Graph Databases. Graph pattern matching, a fundamental problem in graph query processing, has been extensively studied [5]. In sequential settings, Ullmann’s backtracking algorithm [42] has been optimized using various techniques, such as tree indexing [38], symmetry breaking [17], and compression [7]. Join-based algorithms have been developed for distributed environments. These algorithms use cost estimation to optimize join order, with binary-join algorithms [24, 25] estimating costs using random graph models and worst-case-optimal join algorithms [4] ensuring a worst-case upper bound on the cost. Hybrid approaches [31, 43] adaptively select between binary and worst-case optimal joins based on the lower cost. Recent studies have focused on improving cost estimation in graph pattern matching, including decomposing graphs into star-shaped subgraphs [32] and comparing different cardinality estimation methods [35]. Some optimizers, like GLogS [26], search for the optimal plan by representing edges as binary joins or vertex-expansion subtasks. We follow the join-based methods such as [26, 43] due to their compatibility with the relational context for which RelGo is designed.

Bridging Relational and Graph Models. There is a growing interest in studying the interaction between relational and graph models. DuckPGQ [40, 41] has demonstrated support for SQL/PGQ within the DuckDB [2], utilizing the straightforward, graph-agnostic approach to transform and process pattern matching. Index-based methods, such as GQ-Fast [30] and GrainDB [21], work towards construct graph-like index on relational databases to improve the performance of join execution. RelGo leveraged GrainDB’s indexing technique for implementing physical graph operations. In contrast, methods like GRFusion [18] and Gart [39] work towards materializing graph from the relational tables, so that graph queries can be executed directly on the materialized graph. Such methods incur additional storage costs and potential inconsistencies between relational and graph data.

7 CONCLUSIONS

In this paper, we introduce RelGo, a converged relational-graph optimization framework designed for SQL/PGQ queries. We formulate the SPJM query skeleton to better analyze and optimize the relational-graph hybrid queries introduced by SQL/PGQ. After discovering that a graph-agnostic approach can result in a larger search space and suboptimal query plans, we design RelGo to optimize the relational and graph components of SPJM queries using dedicated relational and graph optimization modules, respectively. Additionally, RelGo incorporates optimization rules, such as FilterIntoMatchRule, which optimize queries across the relational and graph components, further enhancing overall query efficiency. We conduct extensive experiments to evaluate the performance of RelGo against graph-agnostic baselines, demonstrating its superior performance and confirming the effectiveness of the proposed optimization techniques. Future work includes integrating RelGo with more widely-used backend databases to expand its applicability and impact.

REFERENCES

- [1] 2024. Apache Age. <https://age.apache.org/>.
- [2] 2024. DuckDB. <https://duckdb.org/>.
- [3] 2024. openCypher. <https://openCypher.org/>.
- [4] Khaled Ammar, Frank McSherry, Semih Salihoglu, and Manas Joglekar. 2018. Distributed Evaluation of Subgraph Queries Using Worst-Case Optimal Low-Memory Dataflows. *Proc. VLDB Endow.* 11, 6 (oct 2018), 691–704. <https://doi.org/10.14778/3184470.3184473>
- [5] Renzo Angles, Marcelo Arenas, Pablo Barceló, Aidan Hogan, Juan Reutter, and Domagoj Vrgoč. 2017. Foundations of Modern Query Languages for Graph Databases. *ACM Comput. Surv.* 50, 5, Article 68 (sep 2017), 40 pages.
- [6] Anonymous. 2024. *Towards a Converged Relational-Graph Optimization Framework (Full Version)*. Retrieved April 17, 2024 from https://www.dropbox.com/scl/fi/pjld1bdkts4j307lh3q0/relgo_paper.pdf?rlkey=65gtap7vix67p4yh9w4vrv196&dl=0
- [7] Fei Bi, Lijun Chang, Xuemin Lin, Lu Qin, and Wenjie Zhang. 2016. Efficient subgraph matching by postponing cartesian products. In *Proceedings of the 2016 International Conference on Management of Data*. 1199–1214.
- [8] Donald D. Chamberlin and Raymond F. Boyce. 1974. SEQUEL: A structured English query language. In *Proceedings of the 1974 ACM SIGFIDET (Now SIGMOD) Workshop on Data Description, Access and Control* (Ann Arbor, Michigan) (SIGFIDET '74). Association for Computing Machinery, New York, NY, USA, 249–264.
- [9] S. Chatterji, S. S. K. Evani, S. Ganguly, and M. D. Yemmanuru. 2002. On the complexity of approximate query optimization. In *Proceedings of the Twenty-First ACM SIGMOD-SIGACT-SIGART Symposium on Principles of Database Systems* (<conf-loc>, <city>Madison</city>, <state>Wisconsin</state>, </conf-loc>) (PODS '02). Association for Computing Machinery, New York, NY, USA, 282–292.
- [10] Surajit Chaudhuri. 1998. An overview of query optimization in relational systems. In *Proceedings of the Seventeenth ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems* (PODS '98). Association for Computing Machinery, New York, NY, USA, 34–43.
- [11] Surajit Chaudhuri and Kyuseok Shim. 1999. Optimization of queries with user-defined predicates. *ACM Trans. Database Syst.* 24, 2 (jun 1999), 177–228.
- [12] Jin Chen, Guanyu Ye, Yan Zhao, Shuncheng Liu, Liwei Deng, Xu Chen, Rui Zhou, and Kai Zheng. 2022. Efficient Join Order Selection Learning with Graph-based Representation. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining* (Washington DC, USA) (KDD '22). Association for Computing Machinery, New York, NY, USA, 97–107.
- [13] Jonathan Goldstein and Per-Åke Larson. 2001. Optimizing queries using materialized views: a practical, scalable solution. *SIGMOD Rec.* 30, 2 (may 2001), 331–342.
- [14] Goetz Graefe. 1995. The Cascades Framework for Query Optimization. *IEEE Data Eng. Bull.* 18, 3 (1995), 19–29. <http://sites.computer.org/debull/95SEP-CD.pdf>
- [15] Immanuel Haffner and Jens Dittrich. 2023. Efficiently Computing Join Orders with Heuristic Search. *Proc. ACM Manag. Data* 1, 1, Article 73 (may 2023), 26 pages.
- [16] Immanuel Haffner and Jens Dittrich. 2023. Efficiently Computing Join Orders with Heuristic Search. *Proc. ACM Manag. Data* 1, 1 (2023), 73:1–73:26. <https://doi.org/10.1145/3588927>
- [17] Wook-Shin Han, Jinsoo Lee, and Jeong-Hoon Lee. 2013. Turboiso: Towards Ultrafast and Robust Subgraph Isomorphism Search in Large Graph Databases. In *Proceedings of the 2013 ACM SIGMOD International Conference on Management of Data* (New York, New York, USA) (SIGMOD '13). Association for Computing Machinery, New York, NY, USA, 337–348. <https://doi.org/10.1145/2463676.2465300>
- [18] Mohamed S. Hassan, Tatiana Kuznetsova, Hyun Chai Jeong, Walid G. Aref, and Mohammad Sadoghi. 2018. Extending In-Memory Relational Database Engines with Native Graph Support. In *Proceedings of the 21st International Conference on Extending Database Technology, EDBT 2018, Vienna, Austria, March 26-29, 2018*, Michael H. Böhlen, Reinhard Pichler, Norman May, Erhard Rahm, Shan-Hung Wu, and Katja Hose (Eds.). OpenProceedings.org, 25–36. <https://doi.org/10.5441/002/EDBT.2018.04>
- [19] Toshihide Ibaraki and Tiko Kameda. 1984. On the optimal nesting order for computing N-relational joins. *ACM Trans. Database Syst.* 9, 3 (sep 1984), 482–502.
- [20] Toshihide Ibaraki and Tiko Kameda. 1984. On the Optimal Nesting Order for Computing N-Relational Joins. *ACM Trans. Database Syst.* 9, 3 (1984), 482–502. <https://doi.org/10.1145/1270.1498>
- [21] Guodong Jin and Semih Salihoglu. 2022. Making RDBMSs Efficient on Graph Workloads Through Predefined Joins. *Proc. VLDB Endow.* 15, 5 (2022), 1011–1023. <https://doi.org/10.14778/3510397.3510400>
- [22] Jan Kossmann, Thorsten Papenbrock, and Felix Naumann. 2022. Data dependencies for query optimization: a survey. *VLDB J.* 31, 1 (2022), 1–22. <https://doi.org/10.1007/s00778-021-00676-3>
- [23] Ravi Krishnamurthy, Haran Boral, and Carlo Zaniolo. 1986. Optimization of Nonrecursive Queries. In *VLDB'86 Twelfth International Conference on Very Large Data Bases, August 25-28, 1986, Kyoto, Japan, Proceedings*, Wesley W. Chu, Georges Gardarin, Setsuo Ohsuga, and Yahiko Kambayashi (Eds.). Morgan Kaufmann, 128–137. <http://www.vldb.org/conf/1986/P128.PDF>
- [24] Longbin Lai, Lu Qin, Xuemin Lin, and Lijun Chang. 2015. Scalable subgraph enumeration in mapreduce. *Proceedings of the VLDB Endowment* 8, 10 (2015), 974–985.
- [25] Longbin Lai, Zhu Qing, Zhengyi Yang, Xin Jin, Zhengmin Lai, Ran Wang, Kongzhang Hao, Xuemin Lin, Lu Qin, Wenjie Zhang, Ying Zhang, Zhengping Qian, and Jingren Zhou. 2019. Distributed subgraph matching on timely dataflow. *Proc. VLDB Endow.* 12, 10 (jun 2019), 1099–1112. <https://doi.org/10.14778/3339490.3339494>
- [26] Longbin Lai, Yufan Yang, Zhibin Wang, Yuxuan Liu, Haotian Ma, Sijie Shen, Bingqing Lyu, Xiaoli Zhou, Wenyuan Yu, Zhengping Qian, Chen Tian, Sheng Zhong, Yeh-Ching Chung, and Jingren Zhou. 2023. GLogS: Interactive Graph Pattern Matching Query At Large Scale. In *2023 USENIX Annual Technical Conference, USENIX ATC 2023, Boston, MA, USA, July 10-12, 2023*, Julia Lawall and Dan Williams (Eds.). USENIX Association, 53–69. <https://www.usenix.org/conference/atc23/presentation/lai>
- [27] LDBC Social Network Benchmark. 2022. <https://ldbouncil.org/benchmarks/snb/>. [Online; accessed 20-October-2022].
- [28] Viktor Leis, Andrey Gubichev, Atanas Mirchev, Peter A. Boncz, Alfons Kemper, and Thomas Neumann. 2015. How Good Are Query Optimizers, Really? *Proc. VLDB Endow.* 9, 3 (2015), 204–215. <https://doi.org/10.14778/2850583.2850594>
- [29] Feifei Li, Bin Wu, Ke Yi, and Zhuoyue Zhao. 2016. Wander Join: Online Aggregation via Random Walks. In *Proceedings of the 2016 International Conference on Management of Data, SIGMOD Conference 2016, San Francisco, CA, USA, June 26 - July 01, 2016*, Fatma Özcan, Georgia Koutrika, and Sam Madden (Eds.). ACM, 615–629. <https://doi.org/10.1145/2882903.2915235>
- [30] Chunbin Lin, Benjamin Mandel, Yannis Papakonstantinou, and Matthias Springer. 2016. Fast In-Memory SQL Analytics on Typed Graphs. *Proc. VLDB Endow.* 10, 3 (2016), 265–276. <https://doi.org/10.14778/3021924.3021941>
- [31] Amine Mhedhbi and Semih Salihoglu. 2019. Optimizing subgraph queries by combining binary and worst-case optimal joins. *Proc. VLDB Endow.* 12, 11 (jul 2019), 1692–1704.
- [32] Thomas Neumann and Guido Moerkotte. 2011. Characteristic sets: Accurate cardinality estimation for RDF queries with multiple joins. In *Proceedings of the 27th International Conference on Data Engineering, ICDE 2011, April 11-16, 2011, Hannover, Germany*, Serge Abiteboul, Klemens Böhm, Christoph Koch, and Kian-Lee Tan (Eds.). IEEE Computer Society, 984–994. <https://doi.org/10.1109/ICDE.2011.5767868>
- [33] Hung Q Ngo, Ely Porat, Christopher Ré, and Atri Rudra. 2018. Worst-case optimal join algorithms. *Journal of the ACM (JACM)* 65, 3 (2018), 1–40.
- [34] Oracle. 2023. *Property Graph Queries (SQL/PGQ)*. International Organization for Standardization. Retrieved June, 2023 from <https://www.iso.org/standard/79473.html>
- [35] Yeonsu Park, Seongyun Ko, Sourav S. Bhowmick, Kyoungmin Kim, Kijae Hong, and Wook-Shin Han. 2020. G-CARE: A Framework for Performance Benchmarking of Cardinality Estimation Techniques for Subgraph Matching. In *Proceedings of the 2020 International Conference on Management of Data, SIGMOD Conference 2020, online conference [Portland, OR, USA], June 14-19, 2020*, David Maier, Rachel Pottinger, AnHai Doan, Wang-Chiew Tan, Abdussalam Alawini, and Hung Q. Ngo (Eds.). ACM, 1099–1114. <https://doi.org/10.1145/3318464.3389702>
- [36] Protocol Buffers. 2024. <https://protobuf.dev/overview/>.
- [37] Yuan Qiu, Yilei Wang, Ke Yi, Feifei Li, Bin Wu, and Chaoqun Zhan. 2021. Weighted Distinct Sampling: Cardinality Estimation for SPJ Queries. In *SIGMOD '21: International Conference on Management of Data, Virtual Event, China, June 20-25, 2021*, Guoliang Li, Zhanhuai Li, Stratos Idreos, and Divesh Srivastava (Eds.). ACM, 1465–1477.
- [38] Haichuan Shang, Ying Zhang, Xuemin Lin, and Jeffrey Xu Yu. 2008. Taming Verification Hardness: An Efficient Algorithm for Testing Subgraph Isomorphism. *Proc. VLDB Endow.* 1, 1 (aug 2008), 364–375. <https://doi.org/10.14778/1453856.1453899>
- [39] Sijie Shen, Zihang Yao, Lin Shi, Lei Wang, Longbin Lai, Qian Tao, Li Su, Rong Chen, Wenyuan Yu, Haibo Chen, Binyu Zang, and Jingren Zhou. 2023. Bridging the Gap between Relational OLTP and Graph-based OLAP. In *2023 USENIX Annual Technical Conference (USENIX ATC 23)*. USENIX Association, Boston, MA, 181–196. <https://www.usenix.org/conference/atc23/presentation/shen>
- [40] Daniel ten Wolde, Tavneet Singh, Gábor Szárnyas, and Peter A. Boncz. 2023. DuckPGQ: Efficient Property Graph Queries in an analytical RDBMS. In *13th Conference on Innovative Data Systems Research, CIDR 2023, Amsterdam, The Netherlands, January 8-11, 2023*. www.cidrdb.org. <https://www.cidrdb.org/cidr2023/papers/p66-wolde.pdf>
- [41] Daniel ten Wolde, Gábor Szárnyas, and Peter A. Boncz. 2023. DuckPGQ: Bringing SQL/PGQ to DuckDB. *Proc. VLDB Endow.* 16, 12 (2023), 4034–4037. <https://doi.org/10.14778/3611540.3611614>
- [42] Julian R Ullmann. 1976. An algorithm for subgraph isomorphism. *Journal of the ACM (JACM)* 23, 1 (1976), 31–42.
- [43] Zhengyi Yang, Longbin Lai, Xuemin Lin, Kongzhang Hao, and Wenjie Zhang. 2021. HUGE: An Efficient and Scalable Subgraph Enumeration System. In *SIGMOD '21: International Conference on Management of Data, Virtual Event, China*, 128–137. <http://www.vldb.org/conf/1986/P128.PDF>

June 20-25, 2021, Guoliang Li, Zhanhuai Li, Stratos Idreos, and Divesh Srivastava

(Eds.). ACM, 2049–2062. <https://doi.org/10.1145/3448016.3457237>